

*nix shell intro

How does the shell run other programs?

- Your shell is a program, just like any other
- Your shell runs other programs in custom environments
- Running other programs is not magic
- Your is also a REPL (read, eval, print, loop)
- Your shell is also a scripting language

Running other Programs

1. Find the executable on disk
 - a. `$PATH`
 - b. `ls` vs. `/bin/ls`
2. Pass along any invocation arguments
 - a. split on whitespace
 - b. need to quote strings with whitespace (or escape the chars)
3. Record the exit status
 - a. `0 == success`
 - b. `non-zero == failure`
 - c. stored in `$?`
 - d. exit status determines truthiness for conditions

Argument Expansion


1. File globbing, expand into filepaths that match pattern
 - a. `*` # as a regex: `/*`
 - b. `?` # as a regex: `./`
 - c. `[ABCDEFGG]` # as a regex: `/[A-G]/`
2. Command substitution, replace with output from command
 - a. ``cmd arg1 arg2`` # deprecated syntax
 - b. `$(cmd arg1 arg2)` # preferred syntax
3. Variable expansion
 - a. `foo=bar` # reference variables, like to set, without `$`
 - b. `$foo` # expand variables to their values with `$`
4. Brace Expansion
 - a. `foo-{bar,baz}` -> `foo-bar foo-baz`
 - b. `file{0..3}` -> `file0 file1 file2 file3`

which which

- type vs. which
- alias vs. function vs. builtin
- your shell vs. subshell
- sourcing

Running other programs

- `fork(2)`: create a copy of the running process
 - `execve(2)`: replace the running process space with an executable file
 - `wait(2)`: suspend execution until child process terminates



```
int pid = fork();
```

Create a new process that is a clone of its parent.

```
exec*("program" [, argv, envp]);
```

Overlay the calling process virtual memory with a new program, and transfer control to it.

```
exit(status);
```

Exit with status, destroying the process.

```
int pid = wait*(&status);
```

Wait for exit (or other status change) of a child.

Input, Output, and Errors

The three always open file descriptors

1. Standard In: fd0
 - `STDIN`
2. Standard Out: fd1
 - `STDOUT`
3. Standard Error: fd2
 - `STDERR`

Open file descriptors are copied when a process forks and preserved in the process space after `execve`.

Output Redirection

1. Redirect to a file

- `> file` # file opened for writing as fd1
- `>> file` # file opened for writing as fd1

2. Read from a file

- `< file` # file is opened for reading as fd0

3. Redirect `STDOUT` of `cmd1` to `STDIN` of `cmd2`

- `cmd1 | cmd2` # `cmd2` will be the parent of `cmd1`

The shell will, after forking, open the appropriate file in the requested fashion, close open file descriptor and replace it with a newly created file descriptor, THEN exec

The Shell is not Magic

1. Commands are files, aliases, functions, or built-ins
2. Commands are executed with a list of arguments
3. Commands read from STDIN and write to STDOUT/STDERR
4. Commands inherit their environment, ENV and CWD